



ARCHITECTURE

EDUCATION

Year	Degree/Exam	Institute	CGPA/Marks
2028	B.ARCH	IIT Kharagpur	8.09 / 10
2022	AISSCE (XII) (CBSE)	Delhi Public School, Motihari	89.40%
2020	AISSE (X) (CBSE)	Delhi Public School, Motihari	93.60%

INTERNSHIPS

AI Engineer Intern Fitsol Gurugram, Haryana <i>Engineered an E2E ESG intelligence pipeline for a B2B Climate-Tech SaaS, extracting sustainability disclosures and predicting corporate carbon emissions</i> - Predicted emissions at log-R2 = 0.78 by stacking ensemble ML models atop a HistGradientBoosting baseline (log-R2 = 0.72) across 35+ engineered ESG predictors - Structured 1800+ BRSR & ARs into a sector-wise ESG knowledge base via end-to-end intelligence pipeline using Python, LLMs & schema-constrained JSON - Accelerated decarbonization by normalizing & indexing disclosures across 90+ sustainability metrics via HTML/PDF parsing into a queryable knowledge base - Delivered enterprise emission prediction dashboard spanning 2000+ firms across 40+ sectors, enabling real-time ESG benchmarking across stakeholders	[Dec '25 - May '26]
Applied Machine Learning Intern UnoJobs Noida, Uttar Pradesh <i>Architected scalable filter-then-rank pipeline using SQL filters, vector search, and skill-graph matching to surface top-50 candidates from 1M+ applicants</i> - Reduced average time-to-hire to 7 days via a production-ready filter-then-rank pipeline using SQL filters, cosine similarity, and the skill-graph matching - Engineered an LLM reranking layer that distilled 50 candidates into a curated shortlist of 10-15 through evidence-grounded intelligent relevance scoring - Structured an automated candidate-ingestion pipeline to extract, normalize, validate, and store 1M+ applicant profiles for low-latency semantic retrieval - Stored relational data and 1,536-dimensional embeddings in Supabase (pgvector), enabling efficient SQL + vector search across the full candidate pool	[Jan '26 - Feb '26]
AI Agent Operations Intern Omnidel.ai Kolkata, West Bengal <i>Built client-facing AI workflow automation tools for an AI-native platform gamifying work for India's informal workforce, using AI-first development tools</i> - Engineered an automated Fireflies API pipeline to capture, and process 250+ Google Meet transcripts for downstream task extraction across client workflows - Captured 170 of 179 discussed tasks at 95% (recall) accuracy by architecting an LLM-based extraction pipeline converting transcripts into structured JSON - Devised a fuzzy-matching pipeline with LLM reranking to reconcile extracted tasks against existing entries, eliminating duplicate task creation at scale - Orchestrated a Telegram BotFather approval workflow enabling one-tap human validation before pushing confirmed tasks to Trello via API integrations	[May '26 - Present]

PROJECTS

Swiggy Delivery-Time Prediction End-to-End MLOps Pipeline Self Project <i>End-to-end ML regression pipeline for food delivery time prediction, covering feature engineering, ensemble modeling, and full MLOps productionization</i> - Achieved 3.13-min MAE and R2 of 0.83 using the stacked ensemble trained on 15+ engineered predictive delivery features from 45000+ Swiggy orders - Architected a preprocessing workflow across 20+ raw features using sk-learn ColumnTransformer, PowerTransformer, encoding, and scaling for regression - Productionized the pipeline with 5-tool MLOps stack (FastAPI, Docker, DVC, MLflow, GitHub Actions), enabling end-to-end deployment, and model versioning	[Jun '25 - Aug '25]
Spotify Hybrid Music Recommendation Engine Self Project <i>Hybrid music recommendation engine combining content-based and collaborative filtering for context-aware personalized Spotify recommendations</i> - Developed a hybrid ranking model blending content and collaborative similarity scores across 9K+ users and 30K songs, balancing relevance and diversity - Engineered features using TF-IDF, one-hot encoding, frequency encoding, and feature scaling via sk-learn ColumnTransformer for content-based filtering - Processed listening data using Dask and constructed a 30K x 962K sparse user-item matrix with SciPy CSR for efficient collaborative filtering at scale	[Aug '25 - Oct '25]
RAG-Based YouTube Q&A System using LangChain & Gemini Self Project <i>RAG-based Q&A system enabling semantic retrieval over YouTube transcripts using LangChain, Gemini, and ChromaDB with optimized recursive chunking</i> - Engineered a retrieval pipeline with Recursive Character Text Splitting and vector embeddings over 1,000+ transcript lines, enabling rapid semantic search - Architected a RAG pipeline using LangChain, Gemini API, and ChromaDB with prompt templates and system prompts to improve overall answer relevance - Integrated the pipeline into a Streamlit application enabling fully live interactive Q&A with sub-3s response latency across chunked YouTube transcripts	[Nov '25 - Dec '25]

SKILLS AND EXPERTISE

Languages: C/C++, Python, SQL, JavaScript | **AI/ML:** Machine Learning, Deep Learning, NLP, LLMs, RAG, LangChain, LangGraph, TensorFlow, Scikit-learn, Keras
Libraries & Frameworks: STL, Pandas, NumPy, Matplotlib, Streamlit, Tailwind CSS, Bootstrap, Seaborn | **Databases & Vector Stores:** MySQL, ChromaDB, Supabase
Developer Tools: Git/GitHub, Docker, Selenium, Jupyter Notebook, VS Code, Codex, Claude Code, OpenClaw, Antigravity | **Visualization:** Power BI, Tableau
Statistics: Descriptive/Interferential Statistics, Hypothesis Testing, Probability, Time Series, Regression & Correlation Analysis, A/B Testing, Distributions

COMPETITION/CONFERENCE

First Runner-Up CityScape: Lucknow Urban Mobility Case Competition IIM Lucknow 2000+ Participants <i>Secured 2nd place among 2,000+ participants by proposing an AI-driven multimodal urban mobility framework integrating EV infrastructure and reforms</i> - Engineered a YOLOv8-based traffic monitoring system with Automatic License Plate Recognition (ALPR) for real-time vehicle detection, violation tracking - Designed a dynamic traffic signal optimization system adapting to real-time traffic density to improve intersection throughput across urban corridors - Projected a 70% reduction in illegal parking and a 30% decrease in traffic congestion through AI-powered traffic monitoring and the adaptive signal control	[Jan '25 - Feb '25]
--	----------------------------

ENTREPRENEURIAL EXPERIENCES

Co-Founder Quiet (Privacy-first AI layer for lawyers and therapists) <i>Converted 3 early ICP-fit therapists into paid pilot users at \$19/mo, conducting 40+ customer discovery interviews across 2 profession verticals (therapy & law)</i> - Eliminated third-party breach risk for workflows by engineering a privacy-first AI note-taking tool using WhisperKit & Ollama for fully on-device transcription	[May '26 - Present]
---	----------------------------

COURSEWORK INFORMATION

Academics: Data Analytics in Building Science, Programming and Data Structures, Computer Aided Design and Simulations, Linear Algebra, Complex Analysis
MOOCs: Stanford CS229; 100 Days of Machine Learning, 100 days of Deep Learning (CampusX); Production RAG (LangChain) & Vector DBs (FreeCodeCamp)

EXTRA CURRICULAR ACTIVITIES

Design: Led branding, creative design, and media operations as **Design & Media Head** for Kshitij, IIT Kharagpur's Annual Techno-Management Festival
Elocution: Achieved **1st Position** in Hindi Elocution at Open IIT, IIT Kharagpur, a prestigious annual pan-institute literary competition with **50+ skilled participants**
Architecture: Won **Gold** at Inter-College D24 (24-hour architecture design challenge) **twice**, conducted jointly by IIT Kharagpur, Jadavpur University & IIST Shibpur
Illumination: Collaborated in a team to craft and assemble **5,000+ lamps**, contributing to JCB Hall's **3rd-place finish** at Illumination 2025, IIT KGP's flagship event